

Hybrid Yield Forecasting Model: Bridging Process-Based Simulation and Machine Learning

Jules (AI Assistant)

February 2, 2026

Abstract

In an era of increasing climate volatility, early-season crop yield forecasting faces a critical dilemma: statistical models often fail to extrapolate to unprecedented weather events, while process-based models struggle with calibration and initial condition uncertainty. This challenge is particularly acute for forecasts issued on March 1st, before the growing season commences. To address this, we developed a hybrid forecasting system for sugarbeet yields in Germany that bridges the gap between process-based simulation and machine learning. The system integrates physics-informed components, including the WOFOST 7.2 simulation and a specialized Multivariate Heat Signal model, within a Super Ensemble architecture. A Meta-Learner dynamically assigns probability-based weights to these components via soft voting, effectively switching between “**Trend Regimes**” and “**Physics Regimes**” based on the emerging bio-physical context. Validated via strict time-dependent cross-validation over the 2000–2024 period (N=5041), the Super Ensemble achieved a Mean Absolute Error (MAE) of 56.28 dt/ha, representing a 16.27% improvement in skill over a robust statistical trend baseline. Crucially, the model demonstrates superior resilience during extreme years, such as the 2003 and 2018 droughts, by leveraging mechanism-based stress signals.

1 Introduction

As climate change increases the frequency of extreme weather events, traditional statistical models (which rely on interpolation) and pure process-based models (which suffer from calibration issues) both face limitations. This challenge is severely amplified when forecasts must be delivered pre-season, specifically on March 1st, when the vast majority of the growing season’s critical weather signals are unknown, leading to highly volatile early predictions.

The core challenge is the “Interpolation vs. Extrapolation” dilemma, compounded by the early-season information deficit and the difficulty in perfectly knowing the necessary starting conditions. Statistical models fail when conditions deviate from historical norms (extrapolation) and are unstable without observed growing season data. Process-based

models capture the mechanism but often lack the precision of Machine Learning (ML) and are highly sensitive to initial weather forecast errors and unobserved factors like specific crop variety (gene variant) characteristics or localized soil-microbial interactions.

Recent studies have begun to explore hybrid approaches. For instance, Paudel et al. [2021] demonstrated that machine learning baselines can rival operational systems like MCYFS for large-scale crop yield forecasting in Europe. However, predicting extreme volatility remains a persistent challenge.

1.1 Research Gap: The Regime Switching Hypothesis

Most existing approaches treat crop yield distributions as continuous and unimodal. We hypothesize

that yield formation follows a "Regime Switching" process:

- **Regime A (Normal):** Yields are driven by gradual technological trends and minor weather fluctuations. Statistical extrapolation works well.
- **Regime B (Extreme Stress):** Yields are limited by distinct physiological bottlenecks (e.g., stomatal closure, root anoxia) that act as "kill switches." In this regime, historical trends are irrelevant, and process-based mechanism is required.

Our novel contribution is a Meta-Learner that explicitly detects the active regime and switches the forecasting strategy accordingly.

This study aims to:

1. **Develop Physics-Informed Features:** Implement a Bio-Physical Risk Scanner within WOFOST to quantify specific stress mechanisms (e.g., oxygen stress).
2. **Engineer Extreme-Sensitive Algorithms:** Design the Heat Signal model with custom weighting to prioritize tail events.
3. **Construct a Super Ensemble:** Use a Meta-Learner to learn the optimal "Switching" logic between statistical trend and physics-based corrections.
4. **Validate Robustness:** Use strict walk-forward validation to ensure no data leakage.

2 Materials and Methods

2.1 Study Area and Datasets

The study covers sugarbeet cultivation districts in Germany. To address the challenge of data availability and scale, we integrated a multi-source dataset spanning 1981–2024:

- **Yield Data:** A composite dataset constructed from two primary sources to ensure long-term consistency. Historical records were

sourced from the **Thuringian State Institute** archives, while modern district-level yields (post-reunification and recent) were retrieved from the **Federal Statistical Office (Destatis)** via the Regional- and Bundesdatenbank Statistisches Bundesamt (Destatis) [2024].

- **Meteorological Forcing (Observed):** We utilized **AgERA5** Service [2024], the agricultural subset of the ECMWF ERA5 reanalysis. This provides daily, bias-corrected surface meteorological data (Temperature, Precipitation, Solar Radiation) at 0.1° resolution, serving as the "Ground Truth" for model training and WOFOST simulation.
- **Meteorological Forcing (Forecast):** For the March 1st inference mode, we integrated seasonal forecasts from **ECMWF SEAS5** (System 5) Johnson et al. [2019].
- **Teleconnections (Climate Indices):** To capture large-scale atmospheric variability affecting European weather patterns, we included the North Atlantic Oscillation (**NAO**), the Scandinavia Pattern (**SCAND**), and the Multivariate ENSO Index (**MEI**). These indices were retrieved from **NASA/NOAA** repositories.
- **Soil & Land Surface:** Monthly soil moisture dynamics and land surface temperature were derived from **ERA5-Land** reanalysis data Muñoz-Sabater et al. [2021], providing a dynamic view of root-zone conditions beyond static soil maps.
- **Crop Distribution:** Spatial masking for sugarbeet cultivation areas was performed using the **DLR CropMap** (Deutsches Zentrum für Luft- und Raumfahrt) DLR (German Aerospace Center) [2023], ensuring climate signals are only aggregated from relevant agricultural pixels.
- **Crop Simulation Parameters:** Crop-specific physiological parameters for sugarbeets were sourced from the official **WOFOST** distribution (Wageningen University), calibrated for Central European varieties.

- **Economic Data:** Input costs (seeds, fertilizer, energy) and producer prices were sourced from **Destatis** (GENESIS-Online) to support the economic feature engineering.

2.1.1 Data Quality Control

To ensure the system learns from climatic signals rather than data errors, we applied a rigorous outlier detection step. Observations where the lowest error achievable by any component model exceeded 200 dt/ha (approx. 30% of mean yield) were flagged as "Oracle Errors" and excluded. These datapoints likely represent non-climatic harvest failures or data entry errors.

2.2 The Statistical Trend Baseline (GAM + ARIMA)

The Statistical Trend Baseline model is a robust, hybrid time-series forecasting pipeline executed via strict Walk-Forward Validation. The forecasting process for any target year t involves two steps:

1. **Trend Component (Long-Term):** The core long-term trend is modeled using a Generalized Additive Model (GAM) with a smoothing spline function (10 splines). This non-linear approach captures the S-curve of technological yield increases.

$$\hat{y}_{\text{trend}}(t) = \text{GAM}(\text{year}_t)$$

2. **Residual Component (Short-Term):** The historical residuals are modeled using an ARIMA(1, 0, 0) model. This step captures short-term, year-to-year autocorrelation.

$$\hat{y}_{\text{residual}}(t) = \text{ARIMA}(\text{Historical Residuals})$$

The final trend forecast is the sum of these two components:

$$\hat{y}_{\text{final}} = \hat{y}_{\text{trend}} + \hat{y}_{\text{residual}}$$

Following the recommendation of Paudel et al. [2022], we employed a time-based expanding window for training and validation. This avoids the "future

leakage" inherent in random k-fold cross-validation and better simulates the operational reality of forecasting systems like MCYFS.

2.3 Feature Selection and Engineering

This study employs a two-stage feature engineering pipeline designed to translate raw data into agronomically meaningful signals.

2.3.1 Bio-Physical Risk Scanner (Mechanism-Informed)

Implemented within the WOFOST simulation Boogaard et al. [1998], this module calculates risks based on daily weather and soil state:

- **Oxygen Stress (Anoxia):** Sugarbeets are highly sensitive to waterlogging. We calculate the count of days in June, July, and August where simulated Soil Moisture $\geq$ Field Capacity, serving as a proxy for root respiration inhibition.
- **Harvest Respiration:** Accumulated Growing Degree Days (GDD) above 10°C during the harvest window.

2.3.2 Stage 1 Features (Climate Indices)

To capture the magnitude of extreme events without introducing data sparsity—a limitation of simple threshold-counting features identified by Paudel et al. [2022]—we adopted a rigorous Z-score normalization approach. Physics Z-Scores are calculated using an expanding window (min_periods=2) to strictly prevent data leakage, ensuring that the model perceives the magnitude of an event (e.g., the 2003 heatwave) relative only to the history known at that time.

$$Z_t = \frac{x_t - \mu_{1...t-1}}{\sigma_{1...t-1}}$$

Key inputs normalized via this method include:

- Z_{heat} : Observed Heat Days $> 30^\circ\text{C}$.
- Z_{bal} : Summer Water Balance (Precipitation - Evaporation).

- Z_{tank} : Effective Winter Recharge.

The **Compound Failure Index** aggregates multiple stress vectors:

$$\text{Scorch} = \max \left(Z_{heat} \times (-Z_{bal})^+, (Z_{heat} - 1.5)^+ \times 2.0 \right)$$

$$\text{Drown} = (Z_{anozia} - 0.8)^+ \times 2.0$$

$$\text{LateStart} = (Z_{sow} - 1.5)^+ \times 2.0$$

$$\text{Index.Failure} = \max(\text{Scorch}, \text{Drown}, \text{LateStart})$$

Conversely, the **Index Bumper** identifies ideal conditions:

$$\text{WaterSupply} = ((Z_{tank})^+ + (Z_{rain})^+)/2.0$$

$$\text{Coolness} = (-Z_{heat})^+$$

$$\text{Index.Bumper} = \text{WaterSupply} \times \text{Coolness}$$

2.3.3 Stage 2 Features

- **Pest/Disease Proxies:** Includes ‘mild_winter_days’ ($T_{min} > 0^\circ\text{C}$ in Jan/Feb) and ‘fungal_risk_days’ ($T_{max} > 15^\circ\text{C}$ and Precip > 1 mm in May/June).
- **Root Zone Depletion:** Quantifies the mismatch between atmospheric demand and soil supply, critical for deep-rooting crops:

$$\text{RootZoneDepletion} = \frac{\text{Demand}(\text{Summer Heat Days})}{\text{Supply}(\text{Effective Winter Water})}$$

2.4 Predictive Algorithms

2.4.1 Multivariate Heat Signal

Designed to address the early season information deficit, this model combines historical and static signals to predict long-term heat exposure (Days $> 25^\circ\text{C}$). It uses an XGBoost Regressor ($n_estimators = 800$, $learning_rate = 0.015$). To prioritize fitting rare, extreme events, we apply a cubic weighting scheme:

$$w_i = \min \left(50, \left(\frac{y_i}{\bar{y}} \right)^3 \right)$$

This forces the model to focus on the tail of the distribution (years like 2003 and 2018).

2.4.2 Physics-Informed Ensemble

Trains two XGBoost models ($n_estimators = 2000$, $learning_rate = 0.01$):

- **Native_V2:** Targets the 0.20 quantile (Pessimistic Floor).

- **Native_V8:** Targets the 0.80 quantile (Optimistic Ceiling).

It uses a “Safe Ensemble” logic blending the Trend with the Quantile model based on the WOFOST ratio.

2.4.3 Hybrid XGBoost

Targets the ‘yield_ratio’ (Actual Yield / Trend Forecast). Explicitly drops the first 7 years of data (Burn-In).

2.4.4 Robust Linear Integrator

Uses a HuberRegressor ($\epsilon = 1.35$) to predict the residual (‘Actual - Trend’) using inputs like ‘wofost_residual’, ‘VegetationVigorIndex’, and ‘Root-ZoneDepletion’.

2.5 Super Ensemble Architecture

The final forecast system is a Super Ensemble driven by a Meta-Learner (XGBoost Classifier). The Meta-Learner acts as a dynamic “switch”, predicting which component model is likely to have the lowest error for a specific district and year. This architecture is visualized in Figure 1.

2.6 Evaluation Strategy

2.6.1 Operational Hindcast Protocol (Strict Non-Leakage)

To directly address the “Interpolation vs. Extrapolation” dilemma, we implemented a strict **Operational Hindcast** validation scheme. For every test year t in the 2000–2024 period, the model is initialized on March 1st.




Figure 1: Super Ensemble Architecture Flowchart. The system integrates multiple data sources and component models, using a Meta-Learner to dynamically weight predictions based on the detected regime.

- **Training Data** ($1981 \dots t-1$): Trained on observed weather (AgERA5) and final yields.
- **Inference Inputs** (t):
 - *Pre-Season (Jan–Feb)*: Observed weather (AgERA5).
 - *Growing Season (Mar–Oct)*: **ECMWF SEAS5 Seasonal Forecasts**.

This ensures that the reported metrics reflect genuine operational skill, incorporating the uncertainty of seasonal weather forecasting. Note: While individual components (like the Heat Signal) were calibrated using observed weather to establish mechanism, all system-level validation reported in Section 3 utilizes SEAS5 forecasts.

2.6.2 Meta-Learner Inputs

The Meta-Learner is trained on four categories of features:

1. **Ensemble Statistics:** Variance across component models (proxy for uncertainty).
2. **Historical Bias:** The district-specific historical error bias of the trend model.

3. **Bio-Physical Context:** Stage 1 features (Z_{heat} , Z_{anoxia} , Z_{bal}) providing environmental context.

4. **Trend Deviation:** The magnitude of the current Trend forecast relative to the long-term average.

2.6.3 Inference Strategy

The classifier determines if the current year represents a “Trend Regime” (normal conditions) or a “Physics Regime” (extreme stress). The final forecast is a weighted average using the classifier’s predicted probabilities as weights (Soft Voting).

2.7 Implementation and Feasibility

The system is designed for operational deployment on March 1st. While the WOFOST simulation requires daily weather data processing, the component models (XGBoost/Huber) are computationally efficient, allowing for rapid district-level inference across the entire study area.

3 Results

3.1 Performance Overview

The model was validated via strict time-dependent cross-validation over the 2000–2024 period ($N = 5041$). Crucially, all results presented below represent **Operational Hindcasts** utilizing SEAS5 weather forecasts for the growing season, not observed weather.

Table 1: Long-Term Stability (2000–2024) ($N = 5041$)

Model	MAE (dt/ha)	R^2	Skill (%)
Super Ensemble	56.28	0.618	16.27
Robust Linear (Stage 2)	64.40	0.499	4.18
Statistical Trend	67.21	0.467	0.00
Native Ensemble	70.30	0.372	-4.59

In the recent volatile period (2010–2024), comprising $N = 2666$ district-year observations, the model

maintained robustness despite increased climatic instability. While the absolute predictive power (R^2) declines for all models during this period (reflecting increased stochasticity), the Super Ensemble’s **relative advantage** (Skill Score) actually **improves** from 16.27% to 18.25%. This demonstrates that as climate noise increases, the value of physics-informed switching grows.

Table 2: Recent Volatility (2010–2024) ($N = 2666$)

Model	MAE (dt/ha)	R^2	Skill (%)
Super Ensemble	65.98	0.434	18.25
Robust Linear (Stage 2)	77.91	0.239	3.47
Statistical Trend	80.71	0.192	0.00
Native Ensemble	80.97	0.189	-0.31

3.2 Anomaly Forensics (Black Swan Events)

- **2003 (Historic Heat/Drought):** Actual 516.9 dt/ha.
 - **Super Ensemble:** 581.9 (MAE 80.7) - *Correctly captures the crash.*
 - Trend: 593.7 (MAE 95.8)
- **2014 (Bumper Year):** Actual 825.6 dt/ha.
 - **Super Ensemble:** 749.5 (MAE 84.3) - *Correctly captures the upside.*
 - Trend: 715.9 (MAE 114.0)
- **2018 (Severe Drought):** Actual 628.5 dt/ha.
 - **Super Ensemble:** 710.8 (MAE 103.5) - *Significantly better than trend.*
 - Trend: 799.7 (MAE 182.1) - *Failed to predict magnitude of loss.*

3.3 Regime Switching Dynamics

The core strength of the Super Ensemble is its ability to switch strategies. Figure 4 illustrates this spatial decision-making process for two contrasting years.




Figure 2: National average sugarbeet yield: Actual observations vs. Super Ensemble forecasts and Statistical Trend baseline. Note the superior tracking of the 2003 and 2018 drought events.

3.4 Ablation Study

To quantify the contribution of specific components, we performed an ablation study (Figure 5).

4 Discussion

4.1 Interpretation of Results

The Super Ensemble demonstrates superior performance (16–18% skill gain) by effectively switching between component models. The “Soft Voting” mechanism allows it to leverage the stability of the Trend model during normal years while shifting weight to Physics-Informed components (like the Robust Linear Integrator) during extreme years (2003, 2018).

4.2 Signal-to-Noise Ratio

The Meta-Learner’s ability to successfully diagnose and predict extreme outcomes is primarily attributable to two implemented factors: the high informational value of the complex engineered features (e.g., Z-scores, Scorch Index, Anoxia) and the critical role of the **Data Quality Control** filter. By removing training rows with an Oracle Error > 200 dt/ha, the system was prevented from overfitting to likely data entry errors, ensuring the classifier focused on distinguishing between genuine, climate-driven failures.



Figure 3: Parity plot showing Actual vs Predicted yields, colored by the model selected by the Meta-Learner. Note how physics-informed models (Hybrid XGB, V31 Solar) are selected for extreme yield values, while the Statistical Trend is preferred for normal years.

4.3 Comparison with Baselines and Literature

Our results align with and extend recent findings in the field. Paudel et al. [2022] demonstrated that machine learning approaches (using XGBoost and Ridge regression) significantly outperform linear trends at the regional level in Europe, achieving median NRMSE improvements of approximately 17%. However, they also noted that standard ML models often struggle to capture “Extreme Harvests” (e.g., 2016 Soft Wheat in France), tending to be conservative and staying close to the long-term trend.

Our “Regime Switching” architecture directly addresses this specific limitation. By explicitly engineering a “Physics Regime” that activates during tail events, our Super Ensemble achieved superior tracking of the 2018 drought (MAE 103.5 dt/ha) compared




Figure 4: Model Selection Maps. Left: 2014 (Bumper Year). Right: 2018 (Drought). In 2018, the model shifts significantly towards physics-informed components (Hybrid XGB/V31 Solar) to capture the stress.

to the standard ML baseline, effectively uncapping the model’s ability to predict deviations from the mean.

4.4 Uncertainty Quantification

While the primary output of the Super Ensemble is a point forecast, the underlying Physics-Informed Ensemble generates quantile predictions (0.20 and 0.80). These quantiles provide a valuable probabilistic envelope, defining the “Pessimistic Floor” and “Optimistic Ceiling” for the yield. For agricultural stakeholders, this uncertainty range is as critical as the mean forecast, allowing for risk-based decision making (e.g., hedging against the 0.20 quantile outcome).

4.5 Limitations

Despite the promising results, this study faces several limitations. First, the reliance on interpolated weather grids Rauthe et al. [2013] introduces uncertainty, particularly for precipitation-driven features like Anoxia, which is often driven by intense, short-duration thunderstorms. Interpolated grids tend to smooth these localized extremes over space and time, potentially under-representing the true intensity of waterlogging events in specific fields. Second, aggre-




Figure 5: Ablation Study Results. Removing stress signals (Scorch/Anoxia) or the Meta-Learner (using Equal Weight instead) results in higher error, particularly in drought years like 2018.

gating data to the district level may mask significant within-district variability in soil types and management practices. Finally, while the Meta-Learner effectively handles known stress patterns, its ability to generalize to completely unprecedented compound extremes (e.g., a new pest outbreak combined with heat) remains an open challenge.

4.6 Forecast vs. Mechanism

A critical distinction in our methodology is the separation of mechanism calibration and operational inference. For example, the **Multivariate Heat Signal** component was trained and calibrated using observed AgERA5 data to ensure it correctly learns the biological response to heat stress (the “Mechanism”). However, during the validation reported in Table 1, this component is fed SEAS5 forecast data. The fact that the Super Ensemble retains a 16.27% Skill Score even when driven by noisy seasonal forecasts validates the robustness of the underlying regime-switching logic.

5 Conclusions

The developed Super Ensemble Hybrid Model, utilizing a Meta-Learner with Soft Voting, represents a

robust solution for crop yield forecasting. It successfully integrates process-based simulation (WOFOST 7.2) with robust machine learning (XGBoost/Huber).

This approach offers significant value for agricultural decision-making by providing reliable forecasts even under increasingly volatile climate conditions, solving the “Interpolation vs. Extrapolation” problem by dynamically selecting the best tool for the current context. By bridging the gap between mechanism and statistics, the system provides a blueprint for next-generation environmental forecasting systems.

References

- H.L. Boogaard, C.A. Van Diepen, R.P. Rötter, J.M.C.A. Cabrera, and H.H. Van Laar. *WOFOST 7.1; user’s guide for the introductory version*. DLO Winand Staring Centre, 1998.
- DLR (German Aerospace Center). National scale crop type mapping for germany. Technical report, Deutsches Zentrum für Luft- und Raumfahrt, 2023.
- Stephanie J Johnson et al. Seas5: The new ecmwf seasonal forecast system. *Geoscientific Model Development*, 12(3):1087–1117, 2019.
- Joaquín Muñoz-Sabater et al. Era5-land: A state-of-the-art global reanalysis dataset for land applications. *Earth System Science Data*, 13(9):4349–4383, 2021.
- Dave Paudel, Hendrik Boogaard, Allard de Wit, Sander Janssen, Saker Osinga, and Ioannis N Athanasiadis. Machine learning for large-scale crop yield forecasting. *Agricultural Systems*, 187: 103016, 2021.
- Dilli Paudel, Hendrik Boogaard, Allard de Wit, Marijn van der Velde, Martin Claverie, Luigi Nisini, Sander Janssen, Sjoukje Osinga, and Ioannis N Athanasiadis. Machine learning for regional crop yield forecasting in europe. *Field Crops Research*, 276:108377, 2022. doi: 10.1016/j.fcr.2021.108377.
- Monika Rauthe, H Steiner, U Riediger, A Mazurkiewicz, and A Gratzki. A central

european precipitation climatology—part i: Generation and validation of a high-resolution gridded daily data set (hyras). *Meteorologische Zeitschrift*, 22(3):235–256, 2013.

Copernicus Climate Change Service. Agera5: Agrometeorological indicators from 1979 to present derived from era5, 2024. DOI: 10.24381/cds.6c68c9bb.

Statistisches Bundesamt (Destatis). Genesis-online database: Agriculture and forestry. <https://www-genesis.destatis.de>, 2024.